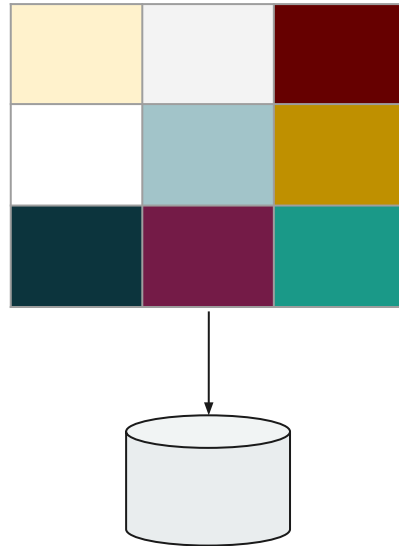




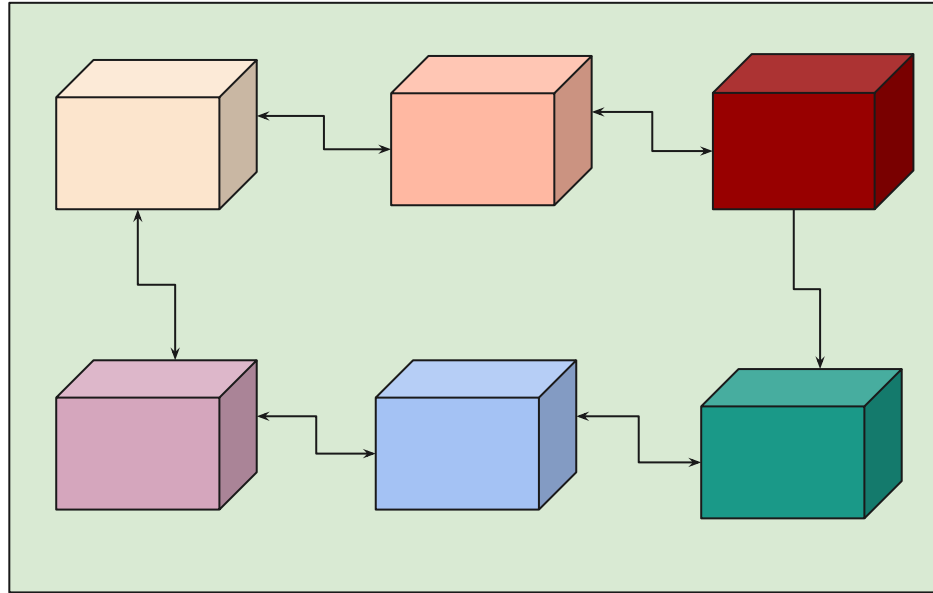
A Complete Guide To gRPC

For Java Developers

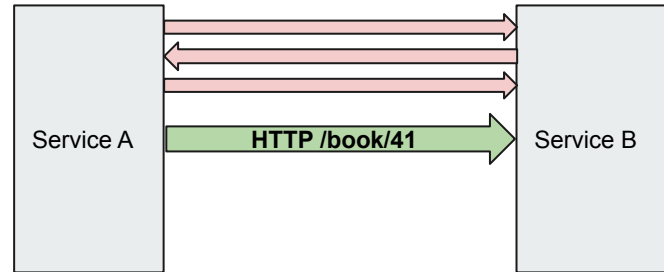
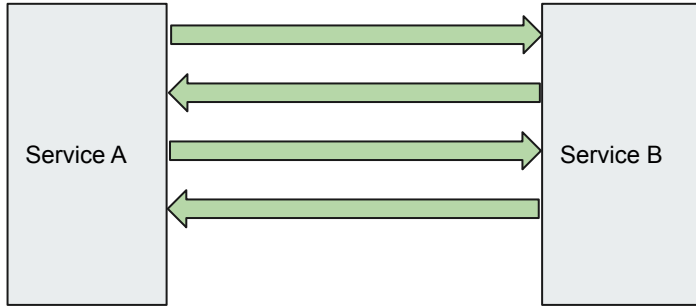
Monolithic Application




MicroServices



Problem - 1: Request & Response Protocol

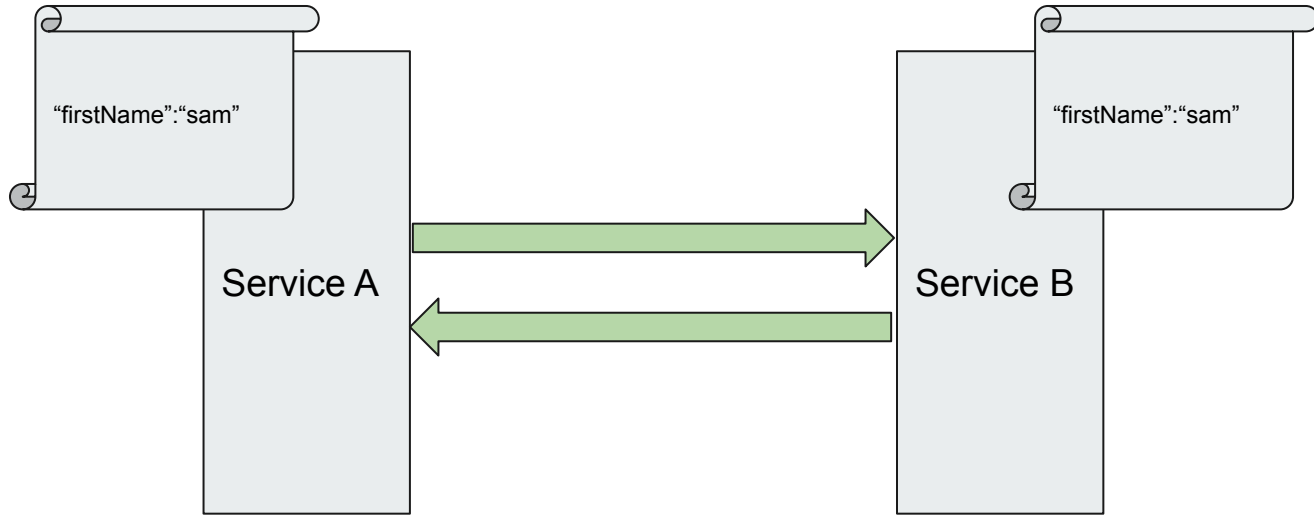


Problem - 2: Headers

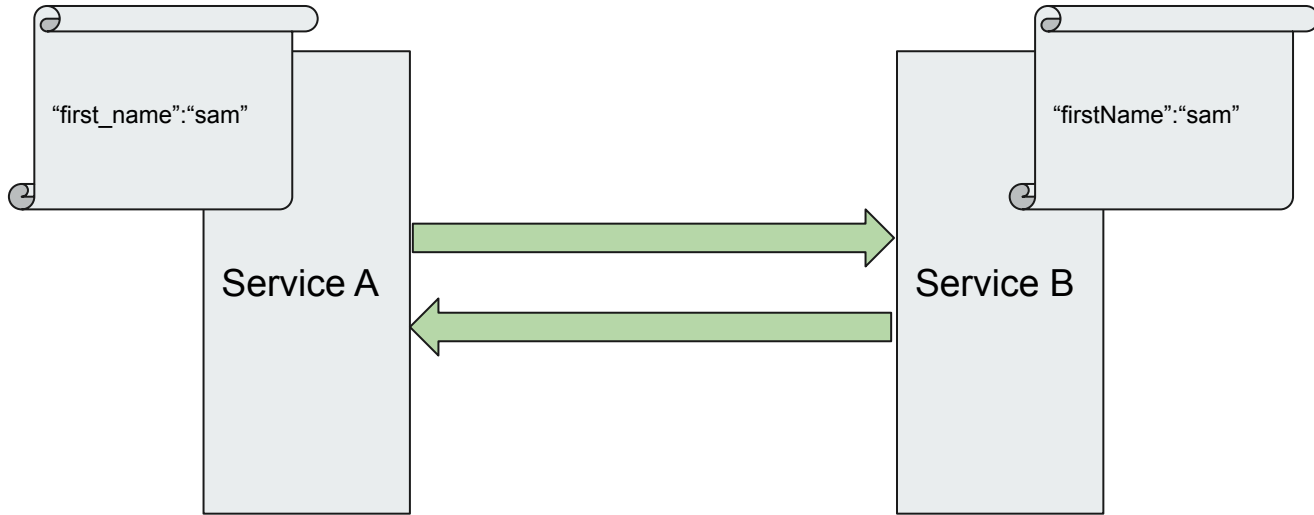


- HTTP is stateless
 - Headers are sent in every request
 - Carries info like Cookie
 - Plain text - relatively large in size
 - Can not be compressed

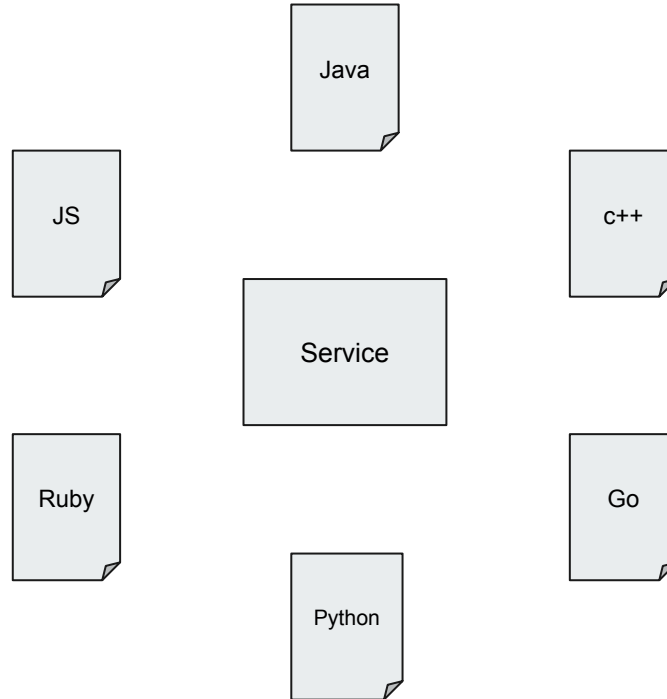
Problem - 3: Serialization & Deserialization



Problem - 4: API Contract

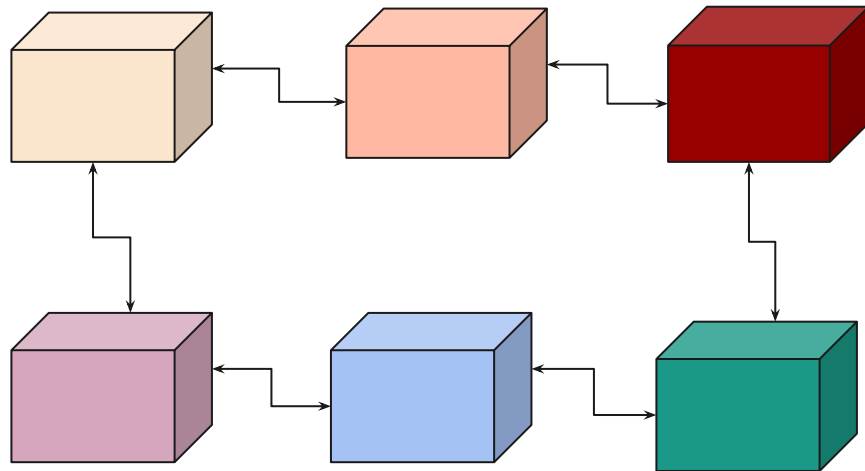


Problem - 5: Client SDK



Stubby

- RPC Framework from Google
- 15 years
- 10 billions reqs / sec!!!!
- Cross-platform
- *Tightly coupled with infrastructure*



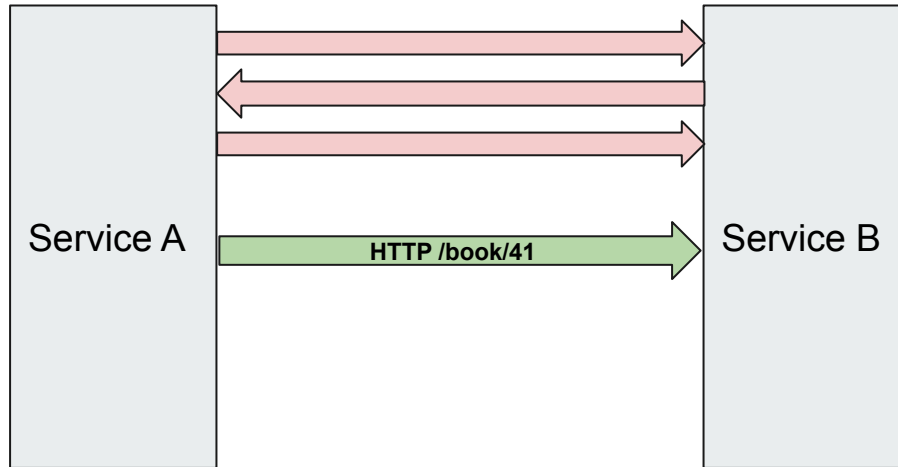
gRPC



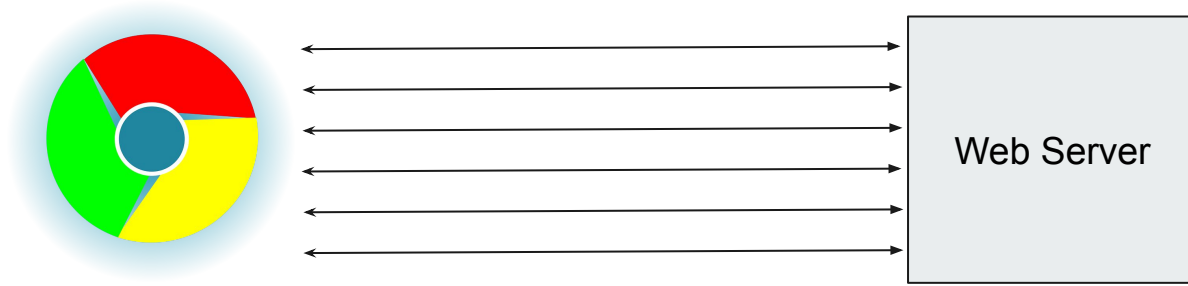
- Developed at Google
- Inspired by Stubby
- Released in 2016
- Adopted by
 - Netflix
 - Microsoft
- Belongs to **CNCF**

HTTP / TCP

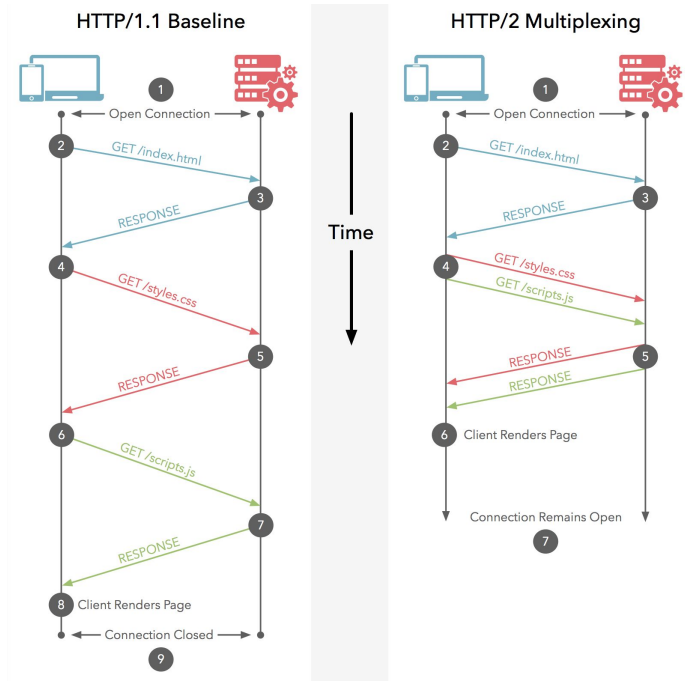
- HTTP 1.1 was released in 1997
- TCP connection - 3 way handshake process
 - Significant amount of time is spent in establishing a connection



HTTP / TCP

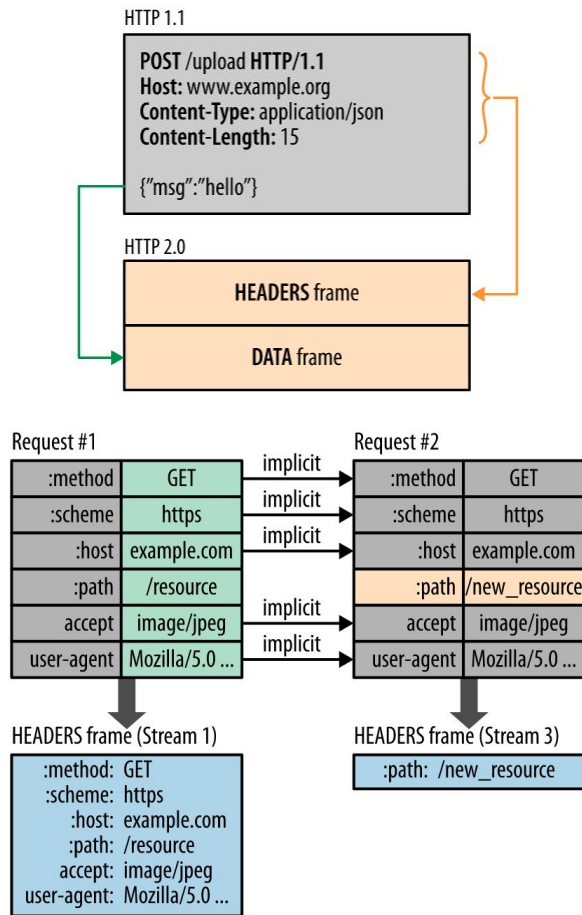


HTTP/1.1 vs HTTP/2



HTTP/2 vs HTTP/1.1

- Binary
- Header Compression
- Flow Control
- Multiplexing



gRPC - Benefits



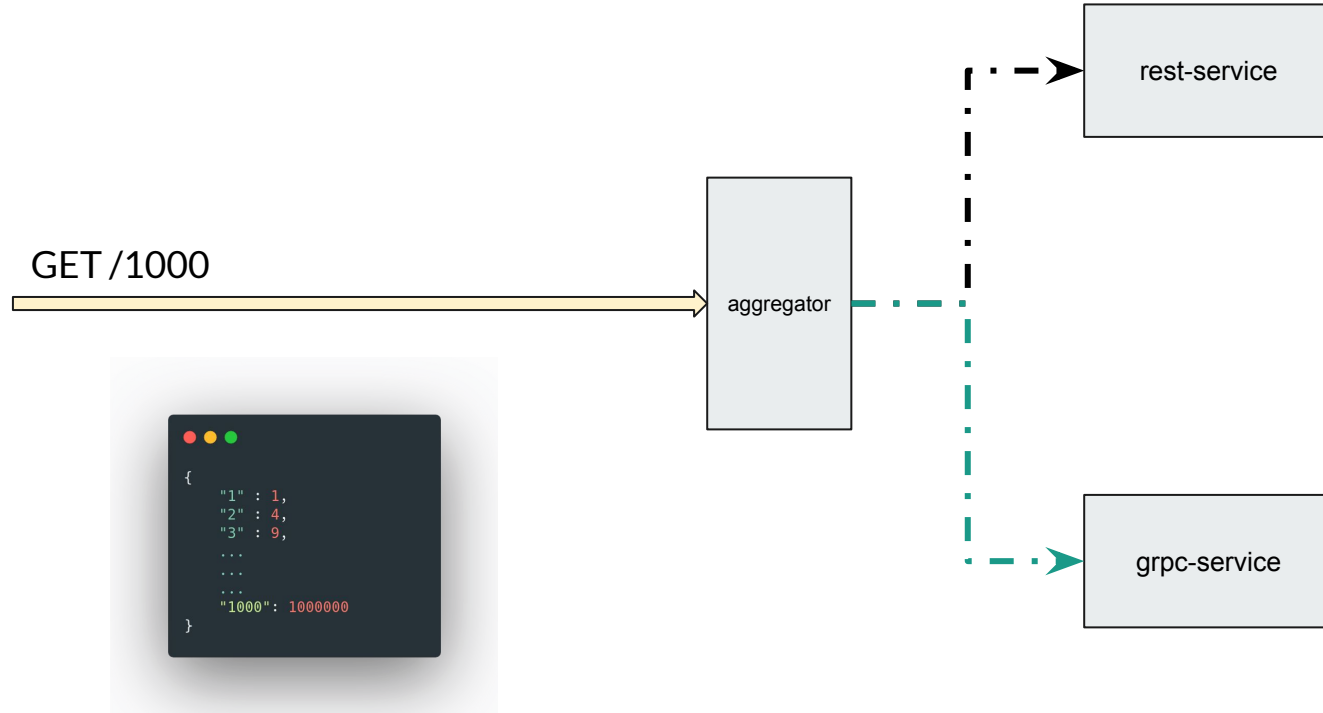
- HTTP2 is default.
 - Binary
 - Multiplexing
 - Flow-control
- Non-blocking, Streaming bindings
- ***Protobuf***
 - Strict Typing
 - DTO
 - Service definitions
 - Language-agnostic
 - Auto-generated bindings for multiple languages
- Great for mobile apps

gRPC vs REST



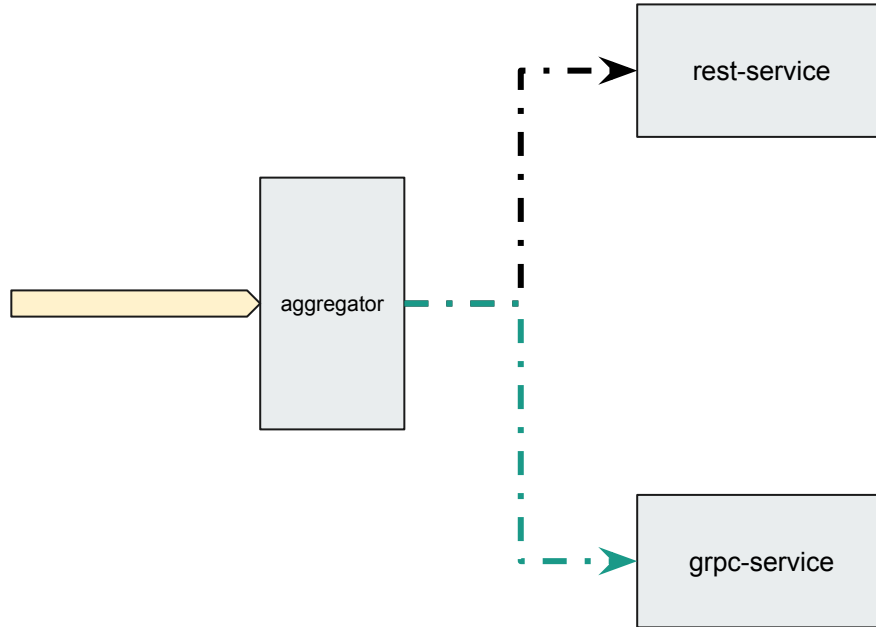
- REST
 - Architectural style
 - Resource oriented
 - ~JSON + HTTP
- gRPC
 - A RPC framework
 - More flexible & Action oriented
 - Specific to inter-services communication

gRPC vs REST



gRPC vs REST

```
// REST  
rest/unary/1000  
  
// gRPC  
grpc/unary/1000  
grpc/stream/1000
```



Course Structure



- Protocol Buffers / protobuf
- gRPC
 - Client/Server application
 - Different types of RPC
 - Client/Server - Game assignment
 - Load Balancing
 - Authentication / Metadata / Contexts
 - Error Handling
 - Deadline
 - Intercepting the requests
- **Bonus:** gRPC - Spring Boot Integration

Note



- Assumption: Experience with Java 8+
 - Pace could be slow or fast depending on your experience :(ul> - **But You would be comfortable with gRPC**
- Would like to show things from scratch



Protocol Buffers

Introduction



- IDL (interface description language) for API
- Platform neutral
- Language neutral
- Serializing/Deserializing structured data
- Very Fast / Optimized for interservices communication
- Provides client libraries automatically for many languages!
 - Java
 - C++
 - Javascript
 - Go
 - Ruby
 - C#
 - Python

Sample Proto

```
public class Person {  
  
    private String name;  
    private int age;  
  
    // getters  
    // setters  
  
}
```

```
message Person {  
    string name = 1;  
    int32 age = 2;  
}
```

Protobuf vs JSON - Performance Comparison



Types




Java Type	Proto Type
int	int32
long	int64
float	float
double	double
boolean	bool
String	string
byte[]	bytes

Proto - Composition

```
message Car {  
  string make = 1;  
  string model = 2;  
  int32 year = 3;  
}
```

```
message Address {  
  int32 postbox = 1;  
  string street = 2;  
  string city = 3;  
}
```

```
message Person {  
  string name = 1;  
  int32 age = 2;  
  Car car = 3;  
  Address address = 4;  
}
```

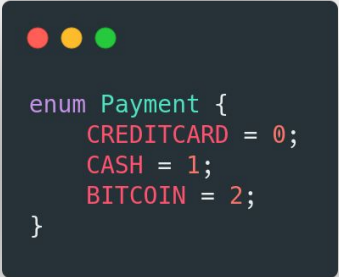


Proto - Collections & Map



Java Type	Proto Type
Collection / List	repeated
Map	map

Proto - Enum



```
enum Payment {  
    CREDITCARD = 0;  
    CASH = 1;  
    BITCOIN = 2;  
}
```

Default Values



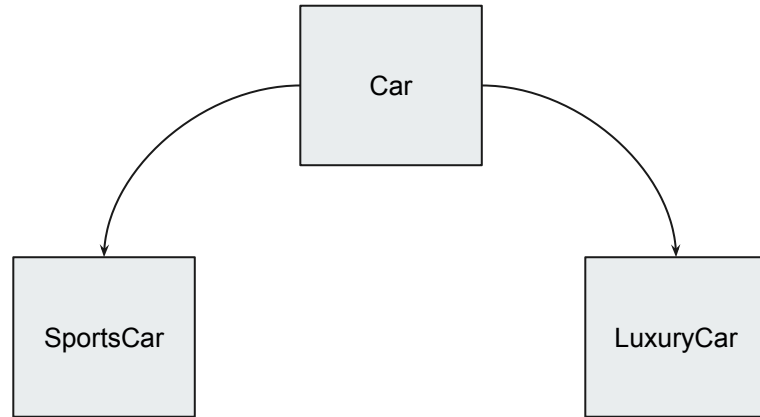
Proto Type	Default
int32 / any number type	0
bool	false
string	empty string
enum	first value
repeated	empty list
map	wrapper / empty map

Proto - Modules / Import



- Proto files can be packaged & imported

Java - Interface / Implementations

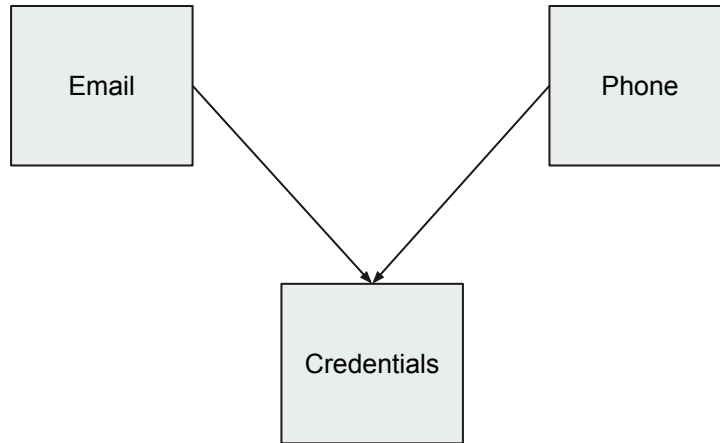


```
drive(new LuxuryCar());
drive(new SportsCar());

//...
//...

public void drive(Car car){
    car.drive();
}
```

Proto - OneOf

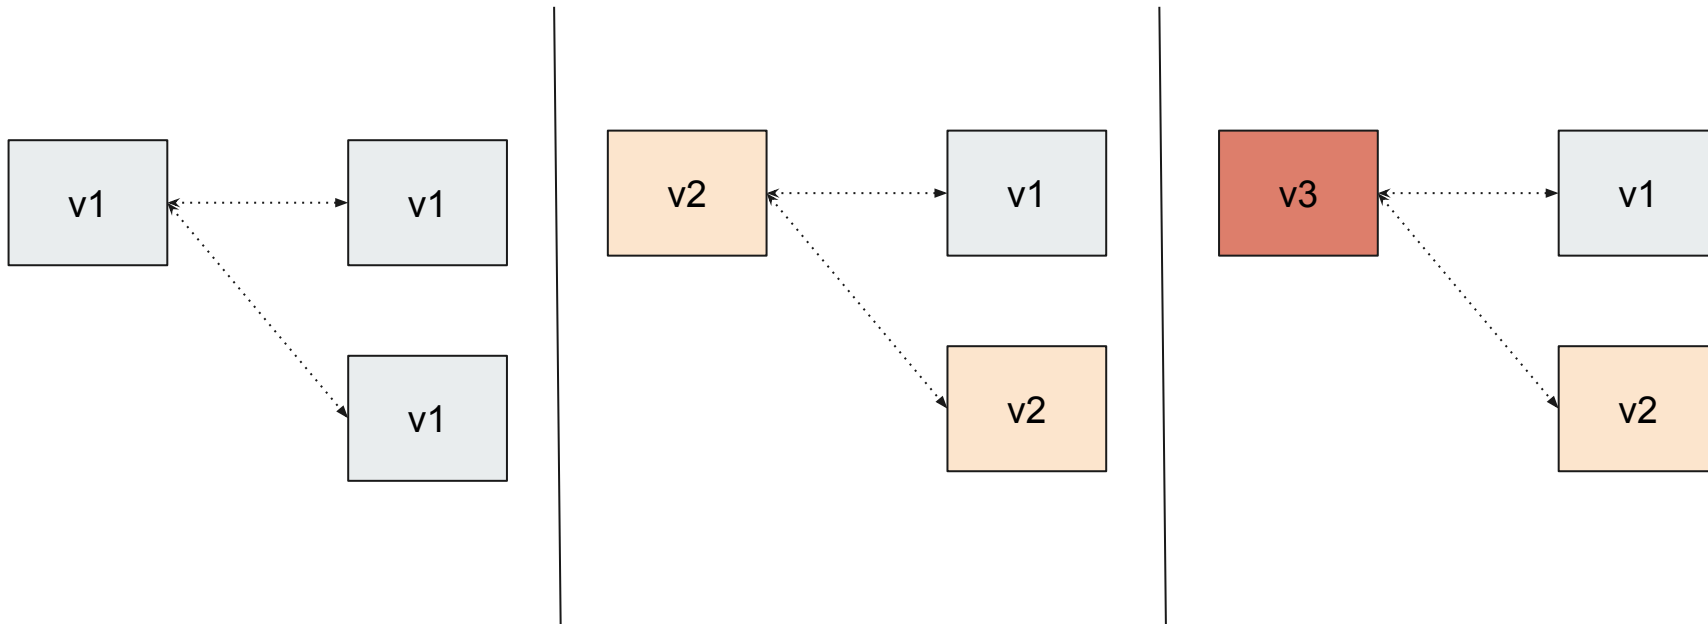


Proto - Field Numbers



- Each Field is assigned with unique number.
- **1-15 for frequently used fields (uses 1 byte)**
- 16-2047 - uses 2 bytes
- 1 is smallest
- $2^{29} - 1$
- 19000-19999 - reserved
- Do not change the field number once it is in use

Proto - Message Format Changes



Proto - Guidelines



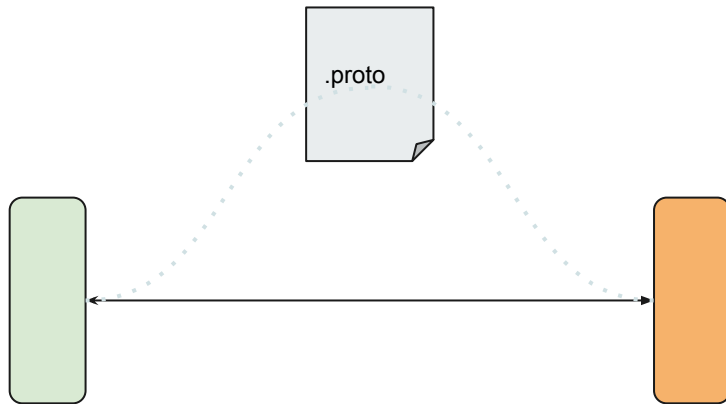
- 1-15 for frequently used fields
- Do not reorder the fields once it is in use
- Adding New fields will not break old proto
- Removing fields will not break old proto
 - Use reserved
- Changing Type
 - int32 → int64 is OK
 - int64 → int32 might be a problem
- Renaming is OK. but be cautious.
- Keep the protos as separate maven-module and add them as dependency in other modules



gRPC

gRPC - Introduction

- High-performance, open-source RPC framework
- Developed at Google
- Client app directly invokes Server method on a different machine
- Service is defined using proto

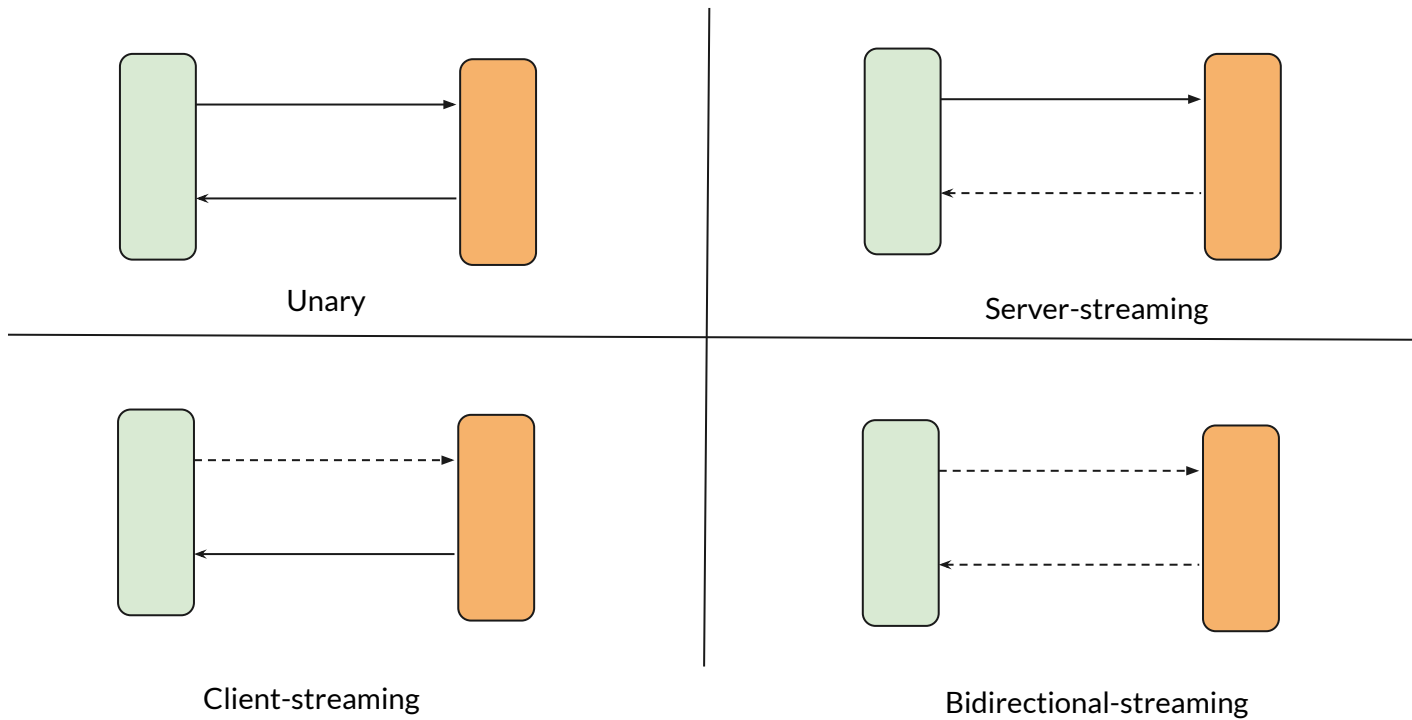


gRPC - Synchronous vs Asynchronous

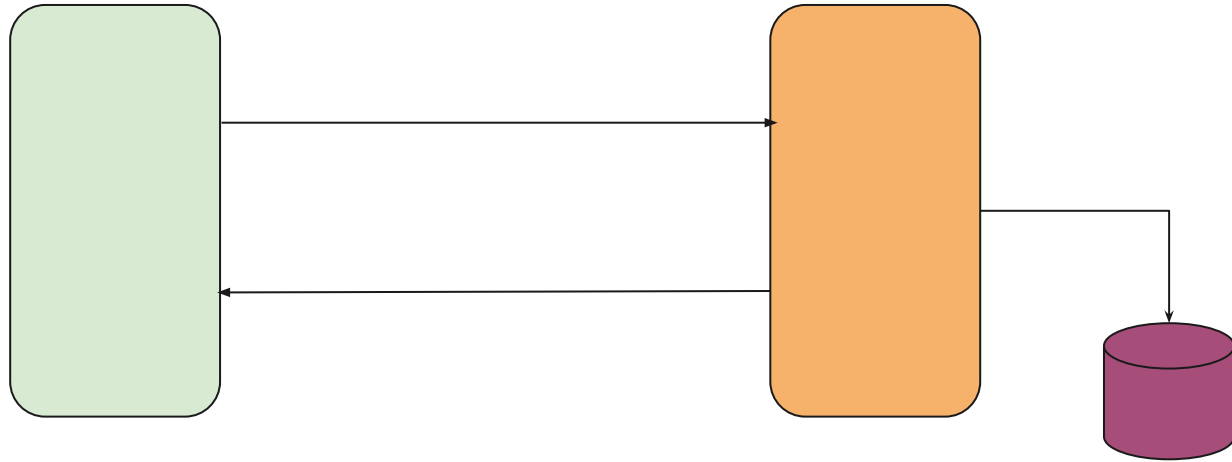


- Client's call to the server can be Sync/Async
 - Sync → blocking / waiting for the response
 - Async → Register a listener for call back
- It is completely up to the client
 - It also depends on the RPC

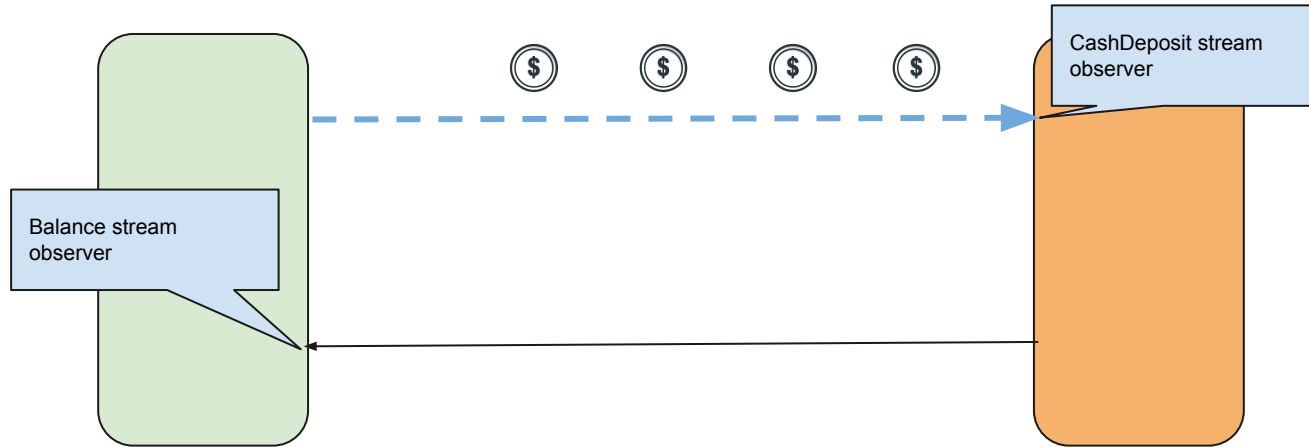
RPC - Types



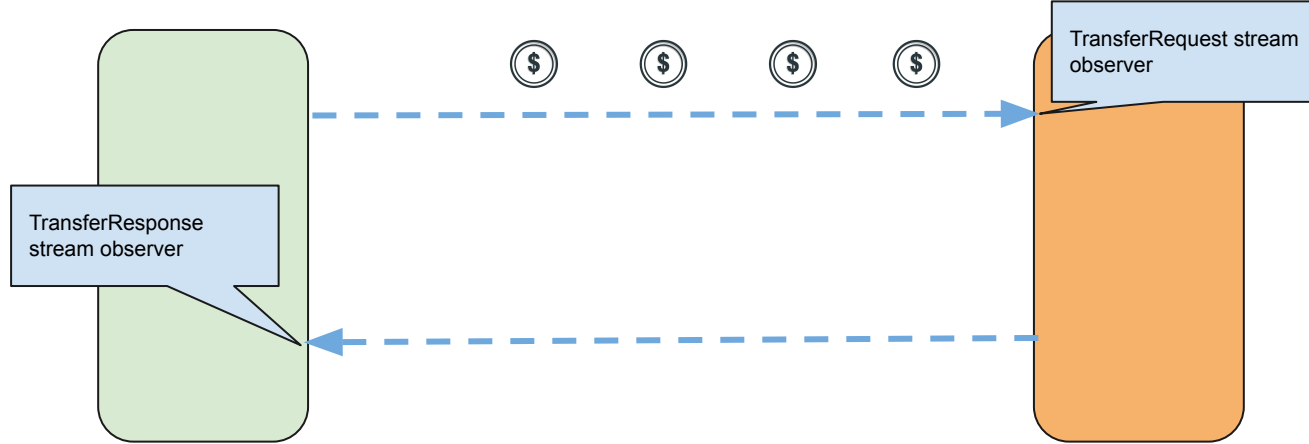
gRPC - Services Communication



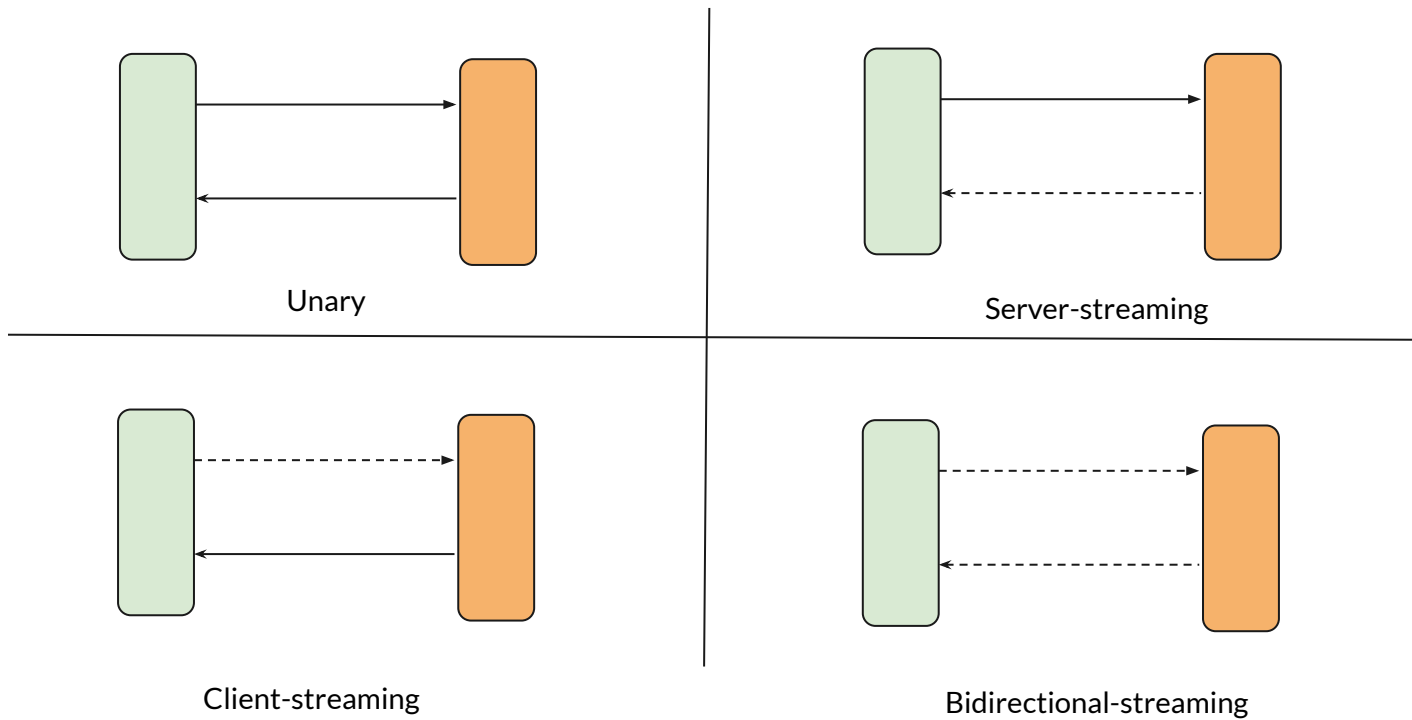
StreamObserver



Bi-directional Stream



RPC - Types



Response - Streaming vs Single



- Streaming
 - Size - Potentially large / unknown
 - Ex: Pagination
 - Uber-eats
 - Receiving side might take too much time to process
 - File upload
 - More efficient than multiple RPC calls
 - Bi-Directional Stream → Client/Server streams are completely independent
- Single
 - More efficient than streaming RPC
 - Size is small

REST - CRUD



GET

POST

```
@RestController
public class BookResource {

    @GetMapping("/book/{id}")
    public Book getBook(@PathVariable int id){
        return bookService.getBookById(id);
    }

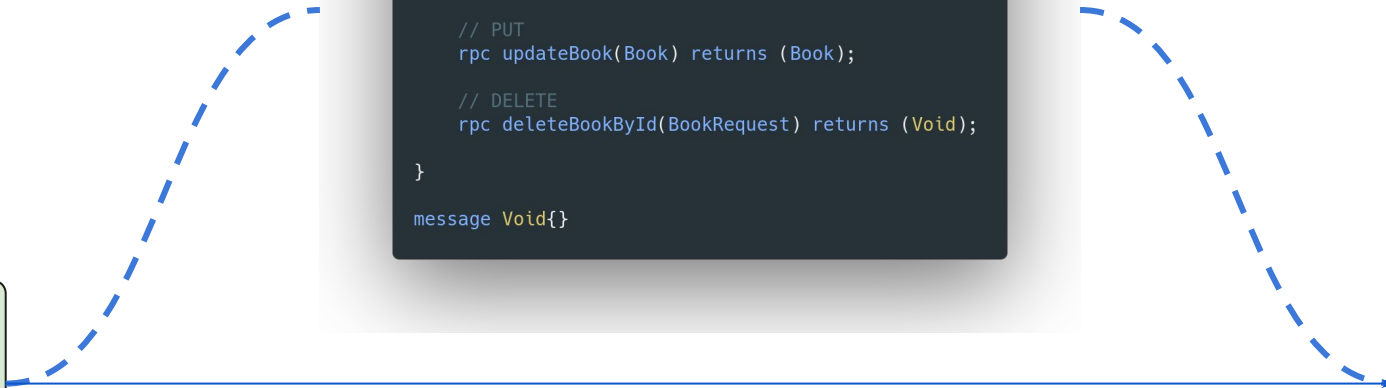
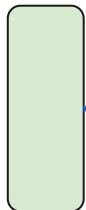
    @PostMapping("/book")
    public Book saveBook(@RequestBody Book book){
        return bookService.saveBook(book);
    }

    @PutMapping("/book")
    public void updateBook(@RequestBody Book book){
        return bookService.updateBook(book);
    }

    @DeleteMapping("/book/{id}")
    public void delete(@PathVariable int id){
        return bookService.deleteBookById(id);
    }
}
```

gRPC - CRUD

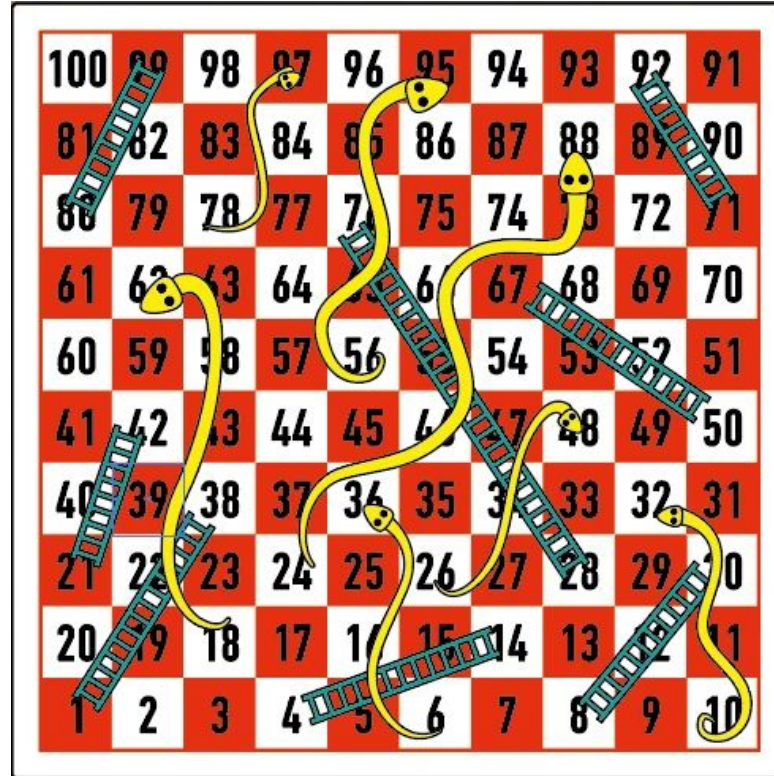
```
service BookService {  
    // GET  
    rpc getBookById(BookRequest) returns (Book);  
  
    // POST  
    rpc saveBook(Book) returns (Book);  
  
    // PUT  
    rpc updateBook(Book) returns (Book);  
  
    // DELETE  
    rpc deleteBookById(BookRequest) returns (Void);  
}  
  
message Void{}
```





Snakes & Ladders

Snakes & Ladders



Hints

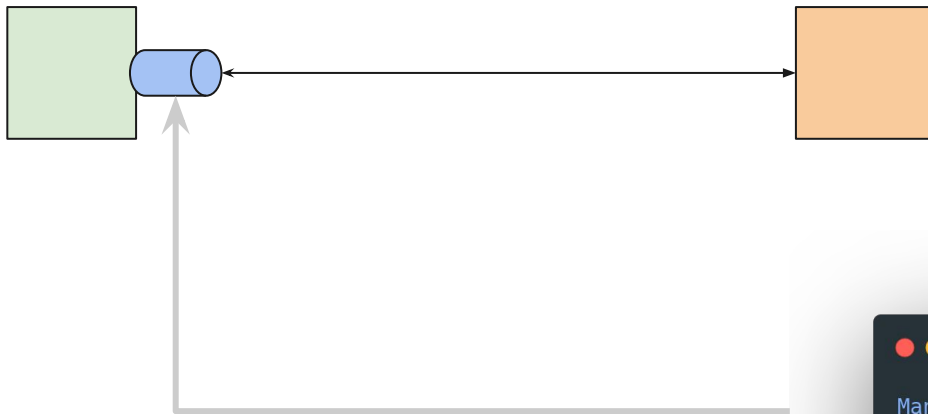


- Client and Server positions start from 0
- Server maintains both client's and server's position
- Client takes the first turn always!
- Client sends its die value to the server
- Server updates the client's position
- Server rolls the die and updates its position
- Server sends the GameState - both client's and server's position once the Server's turn is over or Client wins.
- Client prints the game state as and when it receives the game state
- Game is over when either of them wins!



gRPC - Load Balancing

gRPC - Channel



```
ManagedChannelBuilder.forAddress("localhost", 6565)
    .usePlaintext()
    .build();
```

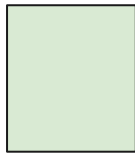
gRPC - Channel



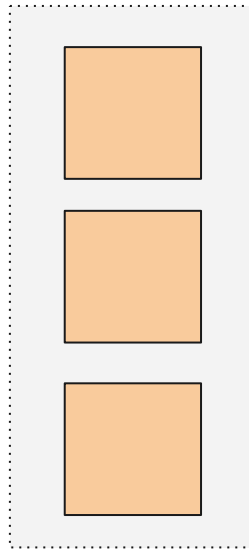
- Channel is an abstract over a connection and represents the connection
- Connection is persistent
- Connection creation is Lazy & established during the first RPC
- 1 channel / connection is enough for client/server communication even for concurrent requests
 - You can also create more channels - but not really required.
 - Channel creation is an expensive process.
- Close when the server shuts down
- Thread safe!
- Can be shared with multiple stubs for the server

gRPC - Load Balancing

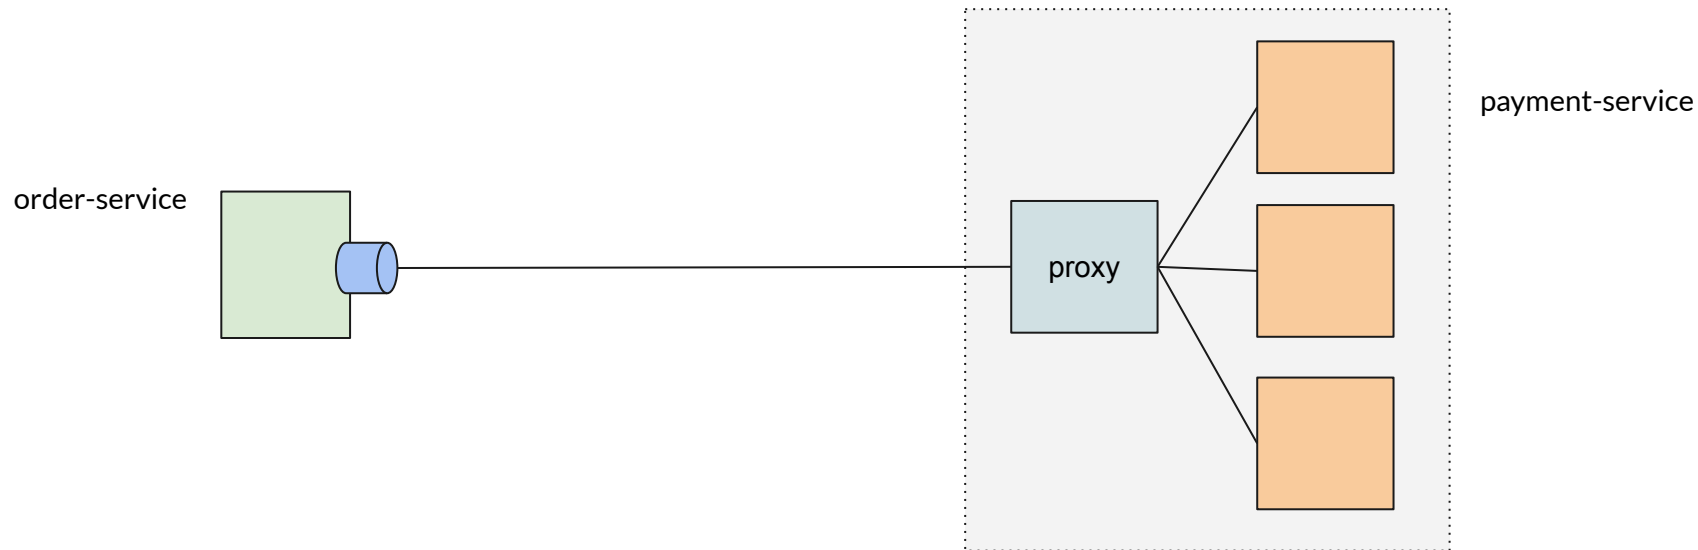
order-service



payment-service



gRPC - Server-Side Load Balancing



nginx

The nginx logo, which consists of a horizontal bar divided into two segments: a teal segment on the left and an orange segment on the right.

```
cd c:\
unzip nginx-1.19.3.zip
cd nginx-1.19.3
start nginx
```

```
// install
brew install nginx

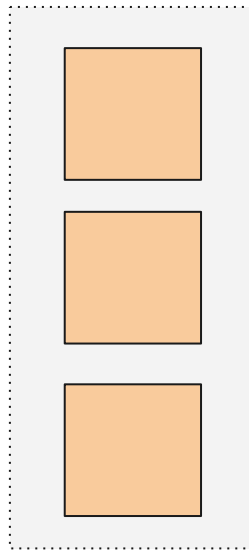
// conf
/usr/local/etc/nginx/nginx.conf
```

gRPC - Load Balancing

order-service

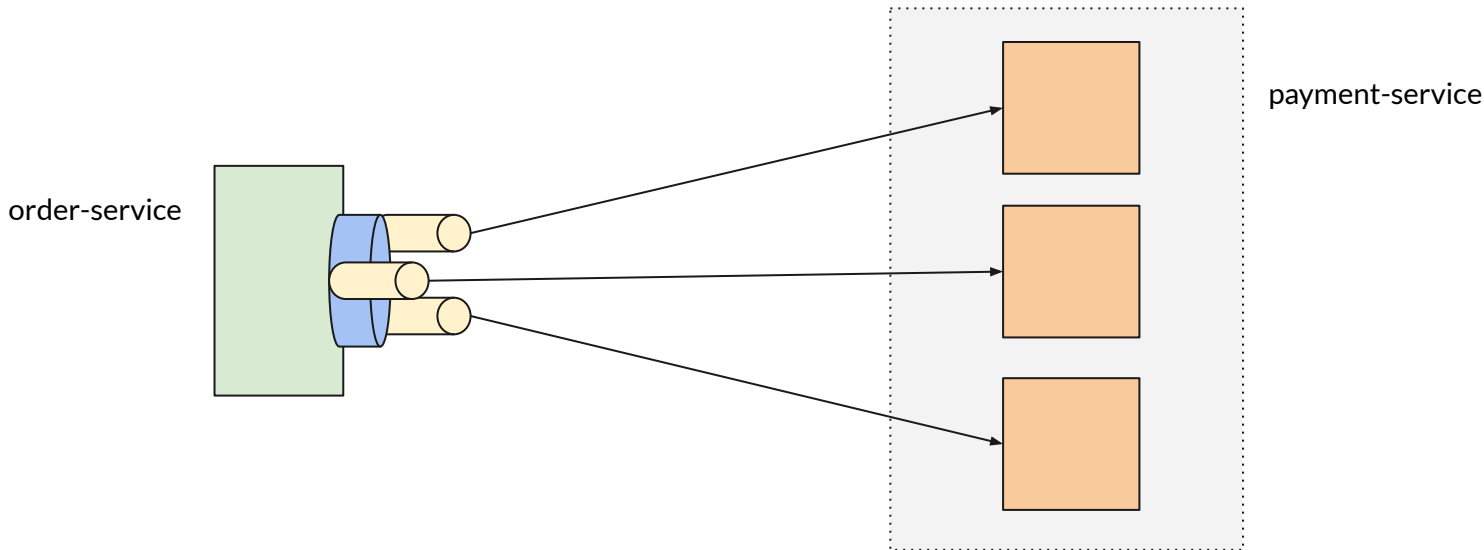


payment-service

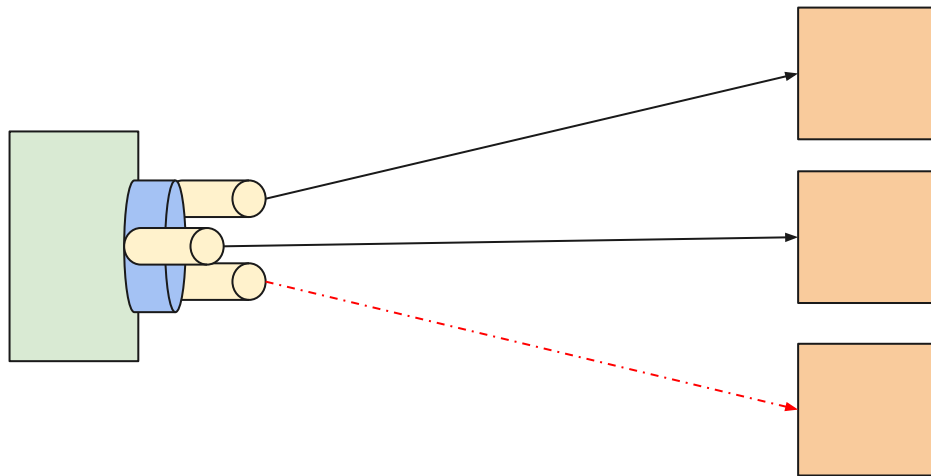


gRPC - SubChannels

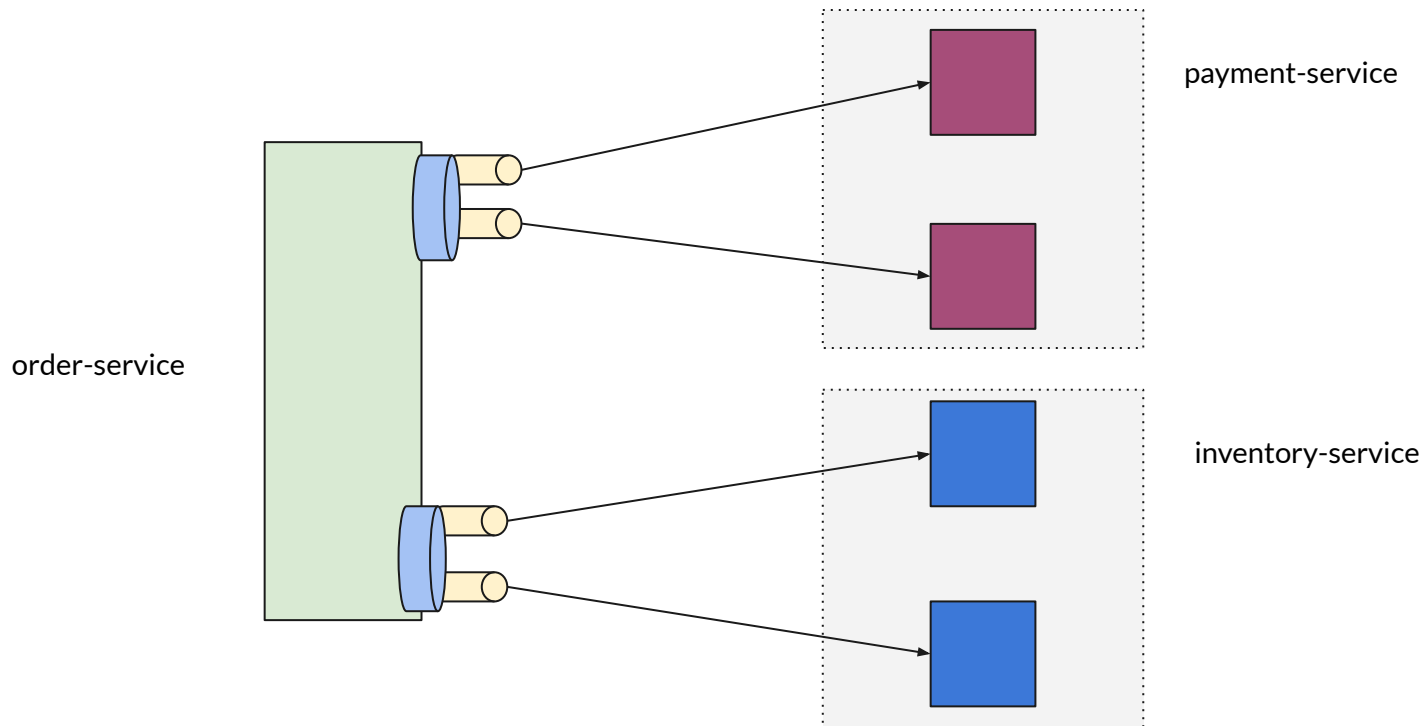
- A Channel can have many subchannels - each subchannel represents a connection to the server.
- Channel chooses subchannel in round robin fashion (*not default*)



gRPC - SubChannels



gRPC - Channels / SubChannels

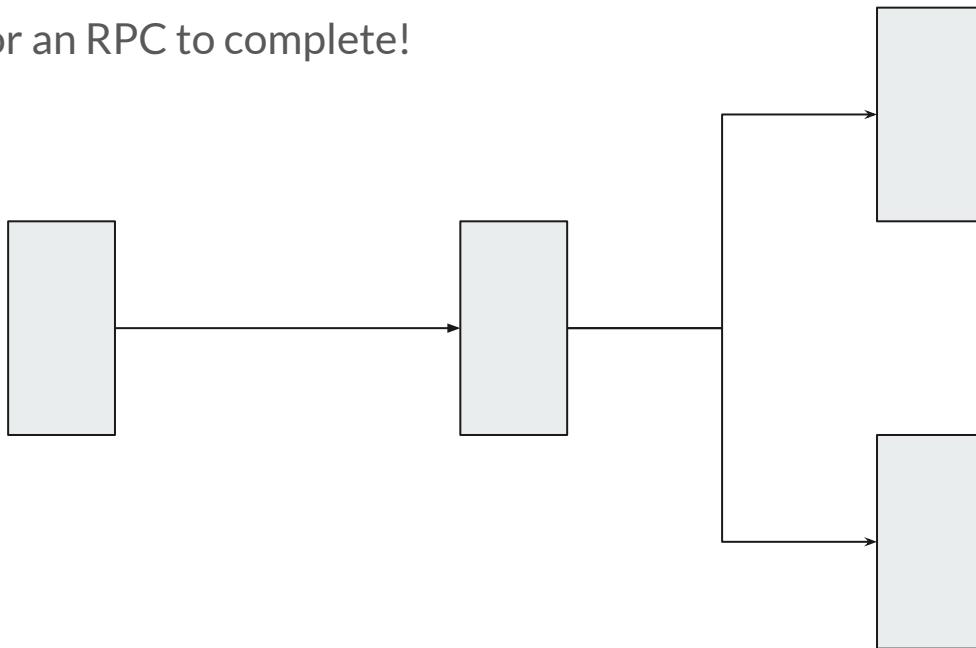




gRPC - Advanced Topics

gRPC - Deadline

- Timeout for an RPC to complete!



gRPC - Interceptor

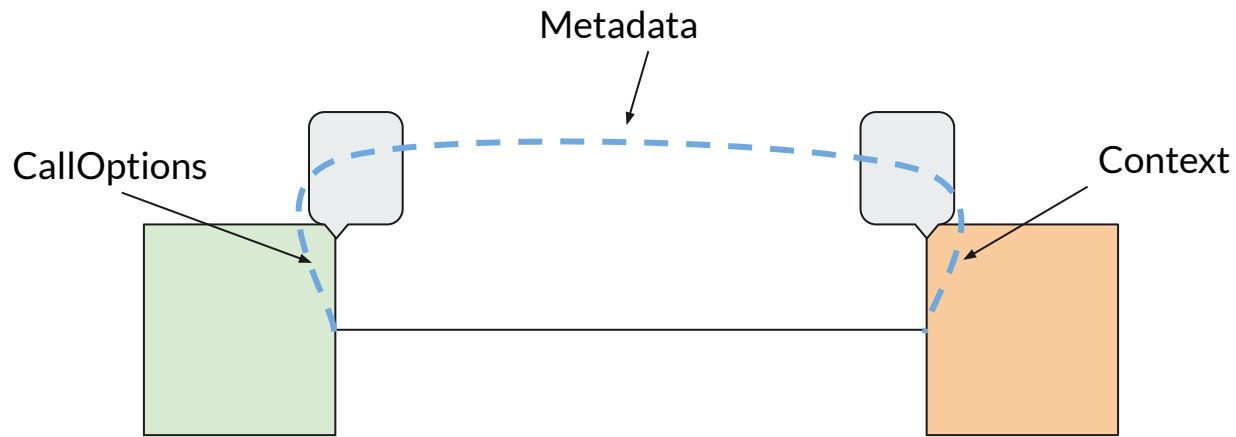


- We can intercept RPCs both at client and server side.
- To handle Cross-cutting Concerns

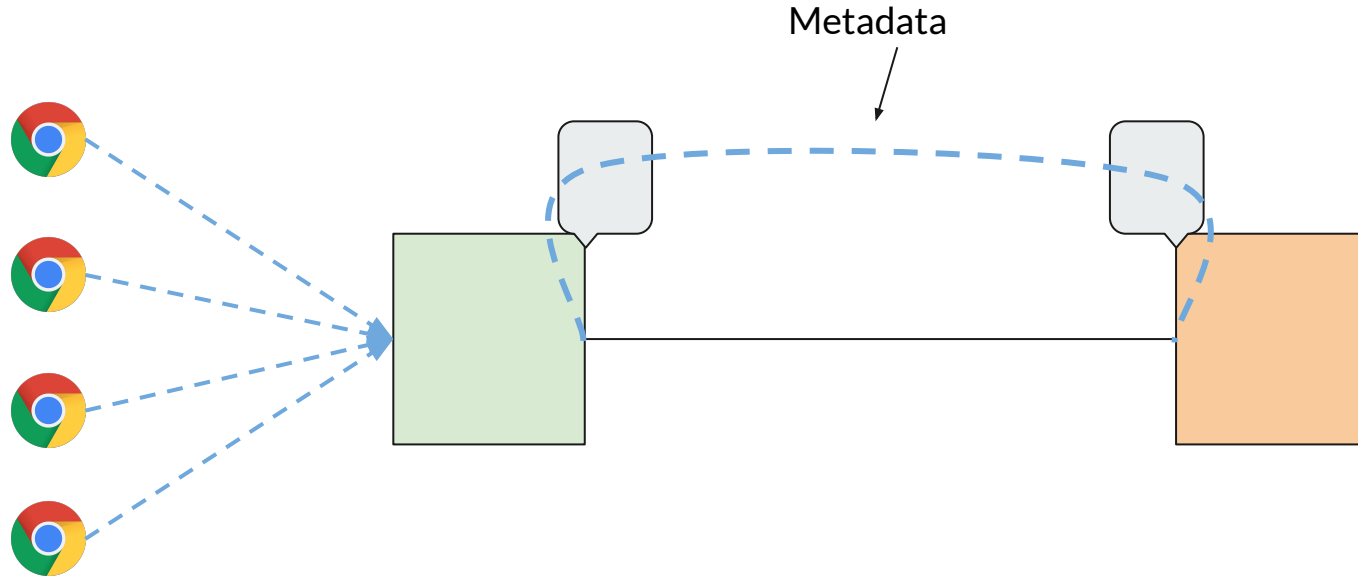
gRPC - Interceptor



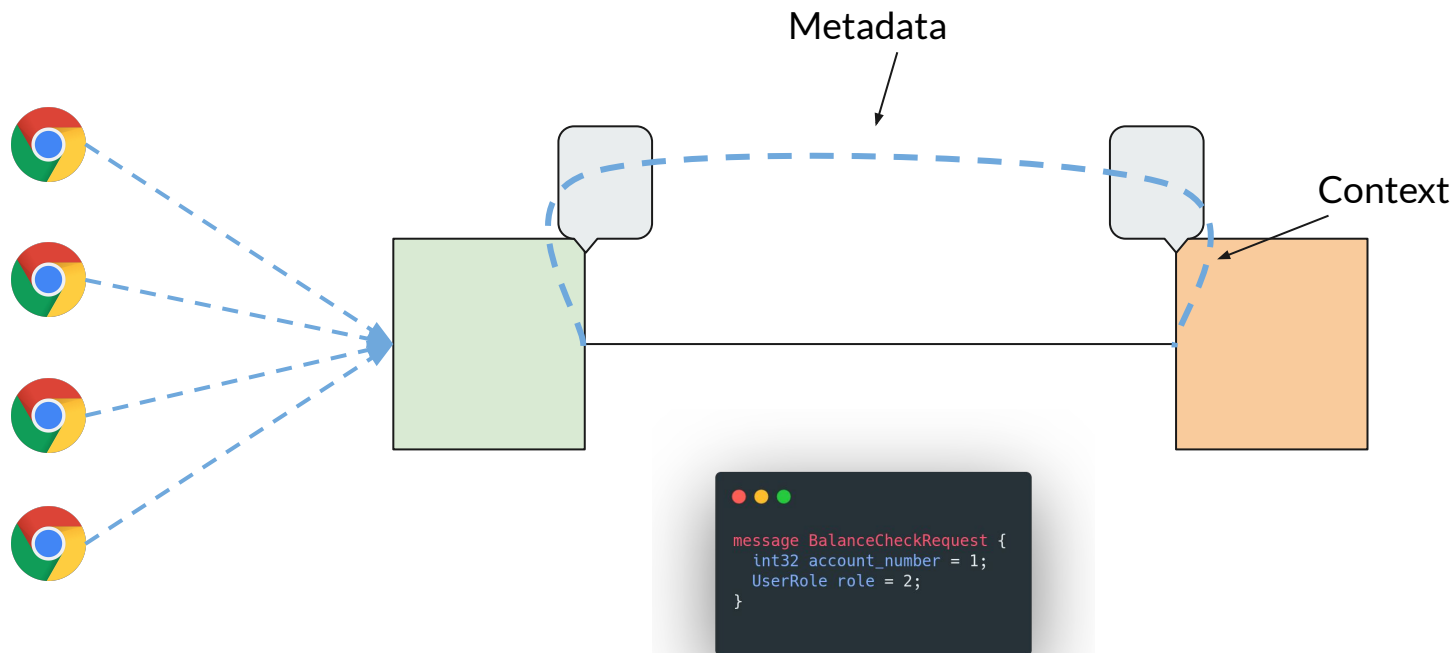
gRPC - Key/Value Pairs



gRPC - Attaching User Token



gRPC - Passing User Role via Context



gRPC - Metadata



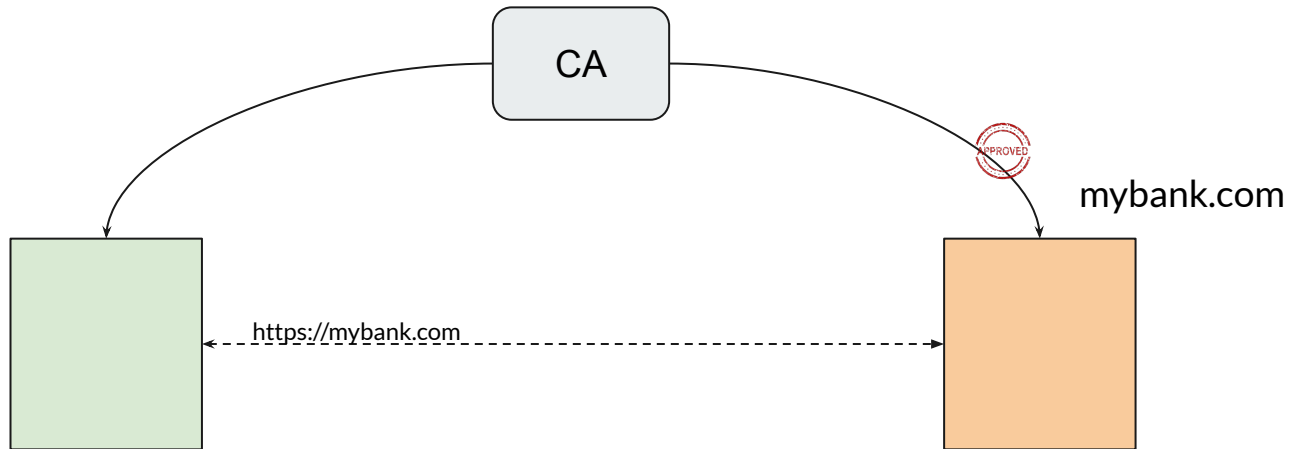
- Key/Value pairs
- Value can be a String or Binary data or an Object
- To pass information between Client & Server

gRPC - Error Handling

- Error channel
 - Status codes
 - Metadata
- Data channel
 - OneOf

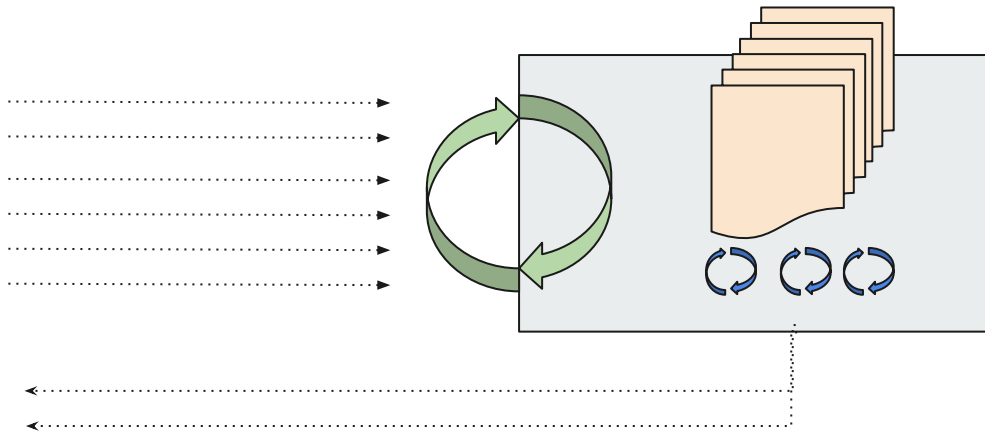
```
service BankService {  
    rpc withdraw(WithdrawRequest) returns(WithdrawResponse);  
}  
  
message WithdrawResponse {  
    oneof response {  
        Money money = 1;  
        Error error = 2;  
    }  
}
```

gRPC - SSL/TLS



gRPC - Executor


- CachedThreadPool - default
- *directExecutor* → Better performance for non-blocking app
- FixedThreadPool → Preferred



gRPC - Executor



```
ServerBuilder.forPort(6565)
    .directExecutor()
    .addService(new GameService())
    .build();
```



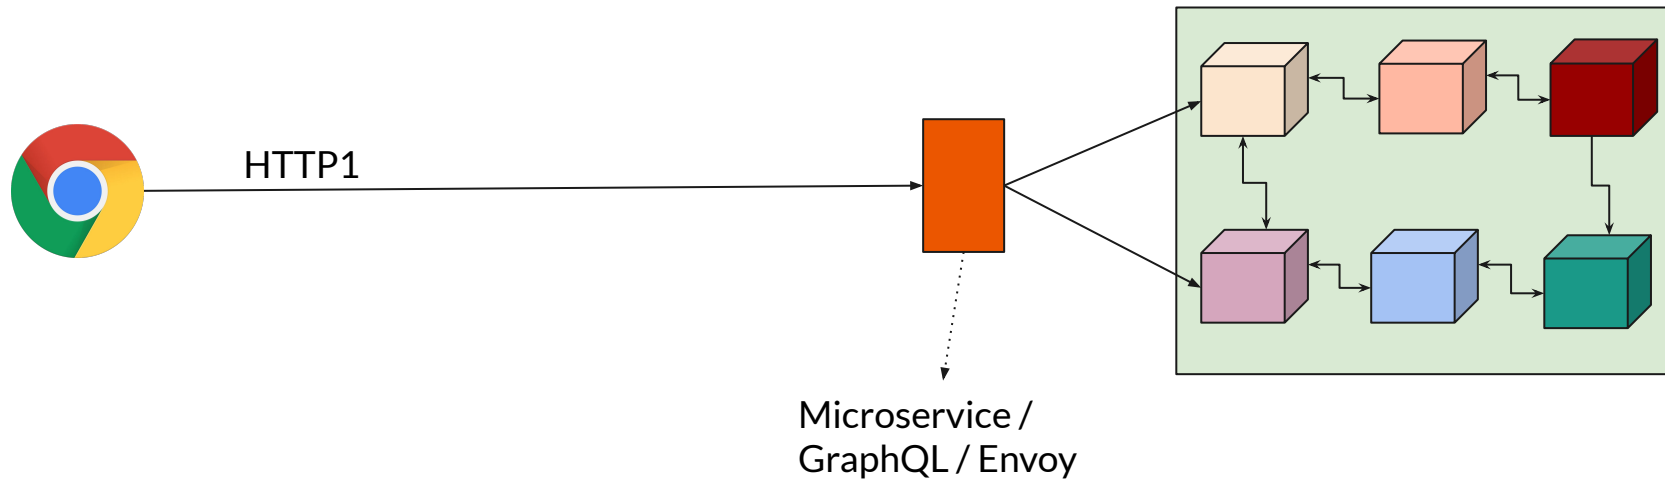
```
ServerBuilder.forPort(6565)
    .executor(Executors.newFixedThreadPool(20))
    .addService(new GameService())
    .build();
```

gRPC - Unit Testing



- InProcessServerBuilder
- InProcessChannelBuilder

gRPC - Web





gRPC - Spring Boot Integration

gRPC - Spring Boot Integration - Server

```
<dependency>
  <groupId>net.devh</groupId>
  <artifactId>grpc-server-spring-boot-starter</artifactId>
  <version>2.9.0.RELEASE</version>
</dependency>
```

```
grpc.server.port=6565
```

```
@GrpcService
public class CalculatorService extends CalculatorServiceGrpc.CalculatorServiceImplBase {

    @Override
    public void getAllDoubles(Input request, StreamObserver<Output> responseObserver) {
        int index = request.getNumber();
        IntStream.rangeClosed(1, index)
            .map(i -> i * 2) // add Thread.sleep to simulate time consuming operation
            .mapToObj(i -> Output.newBuilder().setResult(i).build())
            .forEach(responseObserver::onNext);
        responseObserver.onCompleted();
    }
}
```

gRPC - Spring Boot Integration - Client

```
<dependency>
  <groupId>net.devh</groupId>
  <artifactId>grpc-client-spring-boot-starter</artifactId>
  <version>2.9.0.RELEASE</version>
</dependency>
```

```
grpc:
  client:
    calculator-service:
      address: static://localhost:6565
      negotiationType: plaintext
```

```
@GrpcClient("calculator-service")
private CalculatorServiceGrpc.CalculatorServiceBlockingStub blockingStub;
```

gRPC - Spring Boot Integration

